

36401 Homework 5

```
library(tidyverse)
```

Problem 1

```
library(alr4)
```


```
attach(allshoots)
names(allshoots)
```

```
## [1] "Day" "n" "ybar" "SD" "Type"
```

```
help(allshoots)
```

- a.

```
ggplot(data = allshoots, aes(x = Day, y = ybar, color = as.factor(Type))) +  
  geom_point() +  
  scale_color_discrete(name = "Shoot Length", labels = c("Short", "Long")) +  
  labs(title = "ybar vs Day by Type")
```




While there seems to be a positive linear correlation between Day and \bar{y} for both types, observations with long shoots seem to have a higher slope than those with short shoots.

```
b. out1 = lm(ybar ~ Day * Type)
summary(out1)
```

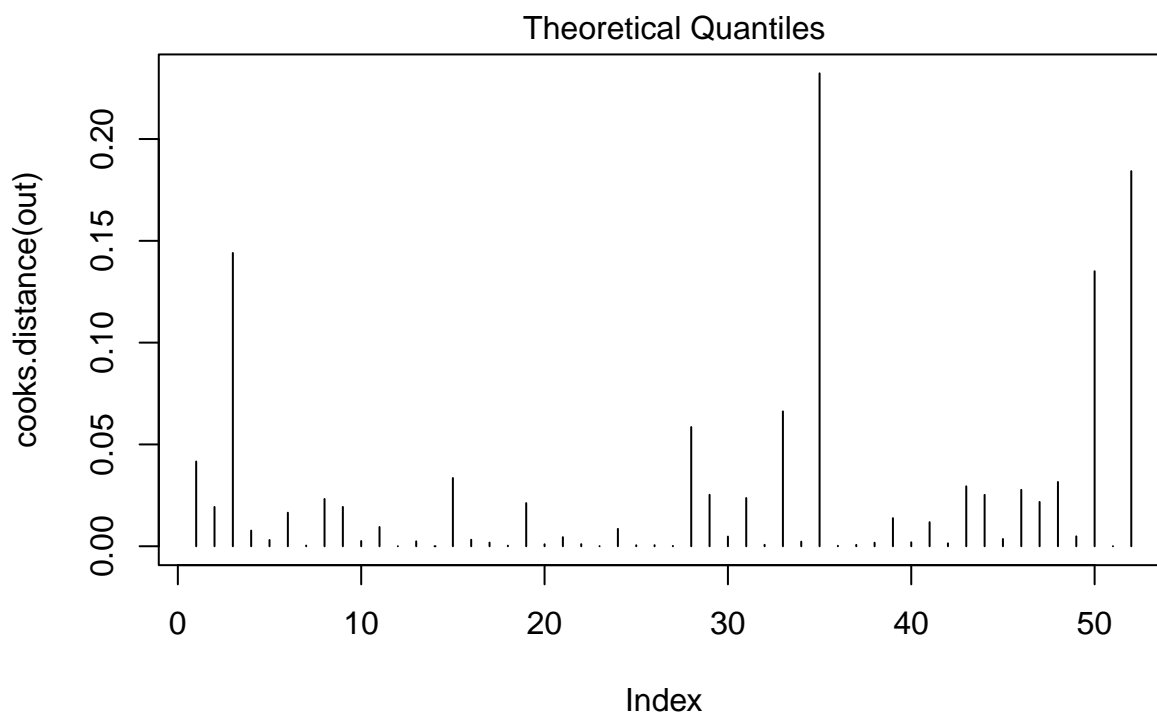

```
##
## Call:
## lm(formula = ybar ~ Day * Type)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.74747 -0.21000  0.08631  0.35212  0.89507
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9.475879   0.230981  41.025  < 2e-16 ***
## Day          0.187238   0.003696  50.655  < 2e-16 ***
## Type         0.339406   0.329997   1.029    0.309
## Day:Type     0.031217   0.005625   5.550 1.21e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5917 on 48 degrees of freedom
## Multiple R-squared:  0.9909, Adjusted R-squared:  0.9903
## F-statistic: 1741 on 3 and 48 DF, p-value: < 2.2e-16
```

The model shows a strong ($R^2 = 99\%$), positive correlation between Day and \bar{y} , with a significant interaction between Day and Type as well.

```
plot_resids = function(out, col = "palevioletred4") {  
  res = rstudent(out)  
  
  #Residual Plot  
  plot(fitted(out), res)  
  abline(h = 0, col = col)  
  
  # QQ Plot  
  qqnorm(res)  
  qqline(res, col = col)  
  
  # Cook's Distance Plot  
  plot(cooks.distance(out), type="h")  
}  
plot_resids(out1)
```



Normal Q-Q Plot



The plot appears to show residuals with approximately equal variance across fitted values of \bar{y} . They also seem to be centered around zero and don't follow any clear pattern. The residuals do not seem to be normally distributed, especially the ones in the lower theoretical quantiles. The Cook's distance plot doesn't appear to show any observations having a too large (>1) influence.

```
library(sandwich)
sandwich_confint = function(out_lm, width){
  alpha = (1 - width) / 2
  V = vcovHC(out_lm)
```

```

se = sqrt(diag(V))
z = -qnorm(alpha/2)
left = out_lm$coef - z*se
right = out_lm$coef + z*se
tmp = cbind(out_lm$coef, se, left, right)
colnames(tmp) = c("Estimate", "se", "left", "right")
print(tmp)
}

sandwich_confint(out1, 0.80)

```

```

##           Estimate          se      left      right
## (Intercept) 9.47587933 0.255944489  9.05488811 9.89687055
## Day         0.18723815 0.003697817  0.18115578 0.19332052
## Type        0.33940568 0.393775435 -0.30829727 0.98710863
## Day:Type    0.03121657 0.006749807  0.02011413 0.04231901

```

For the intercept, the 80% confidence interval is [9.05488811, 9.89687055].

For Day, the 80% confidence interval is [0.18115578, 0.19332052].

For Type, the 80% confidence interval is [-0.30829727, 0.98710863].

For Day:Type, the 80% confidence interval is [0.02011413, 0.04231901].

c.

```

out1_w <- lm(ybar ~ Day * Type, weights = n)
summary(out1_w)

```

```

##
## Call:
## lm(formula = ybar ~ Day * Type, weights = n)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -4.2166 -0.8300  0.1597  0.9882  3.3196
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9.488374   0.238615  39.764 < 2e-16 ***
## Day          0.187258   0.003486  53.722 < 2e-16 ***
## Type         0.485380   0.362496   1.339  0.187
## Day:Type     0.030072   0.005800   5.185 4.28e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.675 on 48 degrees of freedom
## Multiple R-squared:  0.9906, Adjusted R-squared:  0.9901
## F-statistic: 1695 on 3 and 48 DF, p-value: < 2.2e-16

```

```
sandwich_confint(out1_w, 0.95)
```

```
##           Estimate      se      left      right
## (Intercept) 9.48837412 0.213766173  9.00923804 9.96751020
## Day         0.18725802 0.003049236  0.18042346 0.19409259
## Type        0.48537957 0.369744932 -0.34336773 1.31412687
## Day:Type     0.03007228 0.006281888  0.01599204 0.04415253
```

For the intercept, the 80% confidence interval is [9.00923804, 9.96751020].

For Day, the 80% confidence interval is [0.18042346 0.19409259].

For Type, the 80% confidence interval is [-0.34336773, 1.31412687].

For Day:Type, the 80% confidence interval is [0.01599204, 0.04415253].

The confidence intervals are wider for the weighted regression than for the unweighted regression, though this may also have to do with the confidence intervals having a higher confidence level of 95% vs 80%.

Problem 2

```
library(alr4)
attach(BigMac2003)
names(BigMac2003)
```

```
## [1] "BigMac"      "Bread"       "Rice"        "FoodIndex"   "Bus"
## [6] "Apt"         "TeachGI"     "TeachNI"     "TaxRate"     "TeachHours"
```

```
help(BigMac2003)
```

```
a. out2 = lm(FoodIndex ~ BigMac + Bread + Rice + Bus + Apt + TeachGI + TeachNI + TaxRate + TeachHours)
summary(out2)
```

```
##
## Call:
## lm(formula = FoodIndex ~ BigMac + Bread + Rice + Bus + Apt +
##      TeachGI + TeachNI + TaxRate + TeachHours)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -27.0642  -6.3965  -0.0262   5.6928  26.3002
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.09968    11.19872  -0.098   0.9221
## BigMac       -0.20569     0.07798  -2.638   0.0107 *
## Bread         0.44383     0.10564   4.201 9.11e-05 ***
## Rice          0.26881     0.13597   1.977  0.0527 .
## Bus           3.59014     2.83317   1.267  0.2101
## Apt           0.01825     0.00434   4.204 9.02e-05 ***
```


```
## TeachGI      -0.97768    0.86750  -1.127    0.2643
## TeachNI      2.22275    1.13819   1.953    0.0556 .
## TaxRate      0.26530    0.25724   1.031    0.3066
## TeachHours   0.48015    0.20478   2.345    0.0224 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11.86 on 59 degrees of freedom
## Multiple R-squared:  0.7981, Adjusted R-squared:  0.7673
## F-statistic: 25.91 on 9 and 59 DF,  p-value: < 2.2e-16
```

```
plot_resids = function(out, col = "palevioletred4") {
  res = rstudent(out)



  #Residual Plot
  plot(fitted(out), res)
  abline(h = 0, col = col)

  # QQ Plot
  qqnorm(res)
  qqline(res, col = col)

  # Cook's Distance Plot
  plot(cooks.distance(out), type="h")
}
plot_resids(out2, col = "lightpink2")
```



Normal Q-Q Plot



The residuals appear to be random, centered around zero, and have a roughly equal variance across fitted values of `FoodIndex`. However, they don't necessarily seem to follow a normal distribution well, with some of the residuals diverging from the normal line in the Normal Q-Q Plot. None of the observations seem to have a too-large influence on the data based on the Cook's Distance plot.

```
b. sandwich_confint = function(out_lm, width){  
  alpha = (1 - width) / 2
```



```

V = vcovHC(out_lm)
se = sqrt(diag(V))
z = -qnorm(alpha/2)
left = out_lm$coef - z*se
right = out_lm$coef + z*se
tmp = cbind(out_lm$coef, se, left, right)
colnames(tmp) = c("Estimate", "se", "left", "right")
print(tmp)
}
sandwich_confint(out2, 0.99)

```

```

##           Estimate           se           left           right
## (Intercept) -1.0996831 14.538753550 -41.910455301 39.71108903
## BigMac      -0.2056922  0.121622958  -0.547091936  0.13570756
## Bread       0.4438322  0.147985096   0.028433065  0.85923139
## Rice        0.2688081  0.311288199  -0.604988421  1.14260455
## Bus         3.5901381  4.592260279  -9.300491555 16.48076780
## Apt         0.0182453  0.005549614   0.002667345  0.03382325
## TeachGI     -0.9776817  1.044473704  -3.909554657  1.95419126
## TeachNI      2.2227541  1.397504249  -1.700087512  6.14559573
## TaxRate      0.2652971  0.238558093  -0.404343528  0.93493772
## TeachHours   0.4801551  0.253723001  -0.232053936  1.19236413

```

The 99% confidence interval for the **BigMac** covariate is [-0.5471, 0.1357]. This means that 99% of confidence intervals constructed in the same manner as above would contain the true value of **BigMac**. This is the true amount the Food Price Index increases when it takes 1 more minute of labor to purchase a Big Mac. Since the confidence interval contains zero, we can't be sure at the $\alpha = 0.005$ significance level that **BigMac** is a significant predictor for **FoodIndex**.

```

c. out2_bm = lm(FoodIndex ~ BigMac)
# summary(out2_bm)
anova(out2_bm, out2)

```

```

## Analysis of Variance Table
##
## Model 1: FoodIndex ~ BigMac
## Model 2: FoodIndex ~ BigMac + Bread + Rice + Bus + Apt + TeachGI + TeachNI +
##         TaxRate + TeachHours
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      67 27532.9
## 2      59  8299.9  8    19233 17.09 8.026e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

The null hypothesis is that the all of the coefficients other than **BigMac** are zero. Since the p-value is approximately zero, we have significant evidence to reject the null hypothesis. We have sufficient evidence to show that including all of the other covariates as predictors in addition to **BigMac** decreases error.

```

d. pred_error = function(out_lm) {
  return(mean((resid(out_lm))^2 / (1 - (hatvalues(out_lm)))^2))
}

```

```
# Prediction error for all covariates
pred_error(out2)
```

```
## [1] 193.0971
```

```
# Prediction error for big mac-only model
pred_error(out2_bm)
```

```
## [1] 462.8564
```

The prediction error for the model with only BigMac as a covariate is much higher than for the model with all of the covariates included. In conjunction with the F test from earlier, it seems clear that the model with all covariates included gives better predictions.

Problem 3

```
library(mlbench)
data(BreastCancer)
df = BreastCancer[complete.cases(BreastCancer), ]
df$Class = as.numeric(df$Class) - 1
df$Cl.thickness = as.factor(df$Cl.thickness)
df$Bl.cromatin = as.factor(df$Bl.cromatin)
```

```
out3 = glm(Class ~ Cl.thickness + Cell.size + Cell.shape + Marg.adhesion +
           Epith.c.size + Bare.nuclei, data = df, family = "binomial")
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(out3)
```

```
##
## Call:
## glm(formula = Class ~ Cl.thickness + Cell.size + Cell.shape +
##      Marg.adhesion + Epith.c.size + Bare.nuclei, family = "binomial",
##      data = df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.008    0.000    0.000    0.000    1.467
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    34.1527   6186.2156   0.006   0.996
## Cl.thickness.L    68.8230   5753.4447   0.012   0.990
## Cl.thickness.Q    32.7521   3853.4042   0.008   0.993
## Cl.thickness.C    18.7497   3684.1200   0.005   0.996
## Cl.thickness^4    -1.1942   5693.4669   0.000   1.000
## Cl.thickness^5    18.2303   7137.8670   0.003   0.998
```

```

## Cl.thickness^6      1.4076  5862.0001  0.000  1.000
## Cl.thickness^7     -6.8861  3950.9825 -0.002  0.999
## Cl.thickness^8    -35.9175  3708.0154 -0.010  0.992
## Cl.thickness^9     -2.2800  2446.8855 -0.001  0.999
## Cell.size.L        1.8264  44041.0827  0.000  1.000
## Cell.size.Q       -6.7347  22447.2883  0.000  1.000
## Cell.size.C        28.1968  19186.6306  0.001  0.999
## Cell.size^4        33.2267  46489.3609  0.001  0.999
## Cell.size^5        37.2971  56045.7054  0.001  0.999
## Cell.size^6       -2.7392  47892.2153  0.000  1.000
## Cell.size^7        17.0590  30932.5500  0.001  1.000
## Cell.size^8        23.0155  14943.9837  0.002  0.999
## Cell.size^9       -37.0380   8483.8289 -0.004  0.997
## Cell.shape.L       45.0476  43425.9201  0.001  0.999
## Cell.shape.Q      -16.4212  22048.4713 -0.001  0.999
## Cell.shape.C        2.7423  19143.4087  0.000  1.000
## Cell.shape^4       10.0018  45536.4308  0.000  1.000
## Cell.shape^5      -26.8050  54985.5267  0.000  1.000
## Cell.shape^6      -29.6416  46869.3868 -0.001  0.999
## Cell.shape^7      -28.8389  30100.3992 -0.001  0.999
## Cell.shape^8        0.9091  14313.2477  0.000  1.000
## Cell.shape^9     -16.0036   4462.6534 -0.004  0.997
## Marg.adhesion.L    59.3809  12141.6199  0.005  0.996
## Marg.adhesion.Q    27.2403   6991.9637  0.004  0.997
## Marg.adhesion.C   -15.6598   7770.2763 -0.002  0.998
## Marg.adhesion^4   -27.0696  12650.6173 -0.002  0.998
## Marg.adhesion^5   -12.8821  16674.2950 -0.001  0.999
## Marg.adhesion^6   -12.0361  13802.6775 -0.001  0.999
## Marg.adhesion^7    -1.7179   9863.3184  0.000  1.000
## Marg.adhesion^8     8.1805  10220.5200  0.001  0.999
## Marg.adhesion^9     7.2190   6838.1522  0.001  0.999
## Epith.c.size.L     6.7477  19045.6944  0.000  1.000
## Epith.c.size.Q    13.7847   9483.2697  0.001  0.999
## Epith.c.size.C     3.0502   7624.2626  0.000  1.000
## Epith.c.size^4    13.4066  19370.3879  0.001  0.999
## Epith.c.size^5   -20.1535  23772.9094 -0.001  0.999
## Epith.c.size^6    -7.4764  20297.7247  0.000  1.000
## Epith.c.size^7     8.0911  13107.4216  0.001  1.000
## Epith.c.size^8    24.9940   6884.7061  0.004  0.997
## Epith.c.size^9    21.7454   2909.4821  0.007  0.994
## Bare.nuclei2       1.7273     3.9297  0.440  0.660
## Bare.nuclei3      27.5362    19.9990  1.377  0.169
## Bare.nuclei4      32.4423    23.3202  1.391  0.164
## Bare.nuclei5      17.5589    13.7607  1.276  0.202
## Bare.nuclei6      33.7788  27318.7180  0.001  0.999
## Bare.nuclei7      27.5633    19.2736  1.430  0.153
## Bare.nuclei8      -6.1103     6.9944 -0.874  0.382
## Bare.nuclei9      57.4281  19077.2296  0.003  0.998
## Bare.nuclei10     12.0700     8.0537  1.499  0.134
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 884.350  on 682  degrees of freedom
## Residual deviance: 36.473  on 628  degrees of freedom

```

```
## AIC: 146.47
##
## Number of Fisher Scoring iterations: 22
```

The fitted model is shown above.

```
p = predict(out3,type="link")
y = df$Class
n = length(y)
yhat = rep(0,n)
yhat[p >= .5] = 1
T = table(y,yhat)
training_error = (T[1,2] + T[2,1])/sum(T)
print(training_error)
```

```
## [1] 0.01756955
```

The proportion of misclassifications is 0.0176.

Problem 4

```
library(mlbench)
data(Ozone)
help(Ozone)
```

```
Ozone = Ozone[complete.cases(Ozone), ]
#Ozone = subset(Ozone.m, select = -c(V2, V3))
Ozone$V1 = as.numeric(Ozone$V1)
Ozone$month = Ozone$V1
Ozone$oz = Ozone$V4
Ozone$pressureh = Ozone$V5
Ozone$wind = Ozone$V6
Ozone$humidity = Ozone$V7
Ozone$tempS = Ozone$V8
Ozone$tempE = Ozone$V9
Ozone$invHeight = Ozone$V10
Ozone$pressureg = Ozone$V11
Ozone$invTemp = Ozone$V12
Ozone$visibility = Ozone$V13
# Ozone.x = select(Ozone, select = -oz)
# Ozone.x = select(Ozone, select = -V1:V13)
# just take named columns
I = c(14, 16:24)
X = data.matrix(Ozone[,I])
Y = data.matrix(Ozone[,15])

n = nrow(X)
fake = rnorm(20*n)
fake = matrix(fake,n,20)
X = cbind(X,fake)
```

```
a. library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```


```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##     expand, pack, unpack
```

```
## Loaded glmnet 4.1-2
```

```
out4 = cv.glmnet(X, Y, alpha = 1)
plot(out4)
```



```
coefs4 = coef(out4, s="lambda.min")
print(coefs4)
```

```
## 31 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s1
```

```
## (Intercept) -17.146166247
```

```
## month      -0.226098431
```

```
## pressureh  .
```

```
## wind       .
```

```
## humidity   0.098872830
```

```
## tempS      0.127823414
```

```
## tempE      0.308953574
```

```
## invHeight      -0.000261416
## pressureg      .
## invTemp        .
## visibility     -0.002346410
##                0.047788463
##                .
##                .
##                0.029818168
##                .
##                .
##                .
##                .
##                .
##                .
##                -0.565161779
##                .
##                .
##                .
##                .
##                -0.067671907
##                .
##                0.327494032
##                .
##                .
```

```
# plot(b, type="h")
# lambda1 = out$lambda.min
# fit = glmnet(X,Y,alpha=1,lambda=lambda1)

# let's predict some new data
# a = max(abs(Ynew))
# newfit = predict(fit,newx=Xnew,xlim=c(-a,a),ylim=c(-a,a))
# plot(Ynew,newfit)
# abline(a=0,b=1)

# lasso_reg(X, Y)
```

The variables in the final model are month, humidity, temperature at Sandburg, temperature at El Monte, and visibility. Since the model started with 29 variables and went up to 30 according to the plot, it seems that the fake variables were in the model first since there were only 10 real variables.

```
b. lasso = function(X,Y){
  tmp = cv.glmnet(X,Y)
  lambdaCV = tmp$lambda.min
  out = glmnet(X,Y,lambda=lambdaCV)
  return(out)
}

lasso_CI = function(X,Y,j){
  ### confidence interval for beta[j]
  ### in high dimensional regression
  Z = X[,-j]
  X = X[,j]
```

```

tmp1 = lasso(Z,Y)
beta = as.matrix(as.numeric(tmp1$beta))
fitted = Z %*% beta
R = Y - fitted
tmp2 = lasso(Z,X)
beta = as.matrix(as.numeric(tmp2$beta))
fitted = Z %*% beta
S = X - fitted
beta.hat = sum(R*S)/sum(S^2)
sigma = sqrt(mean((R - beta.hat*S)^2))
se = sigma/sqrt(sum(X^2))
return(list(beta.hat=beta.hat,se=se))
}
lci = lasso_CI(X, Y, 5)

```

```

z = -qnorm((1 - 0.9)/2)
lower = lci$beta.hat - z * lci$se
upper = lci$beta.hat + z * lci$se
paste("[", lower, ", ", upper, "]", sep = "")

```

```
## [1] "[0.364491614174825, 0.380672146895094]"
```

So the confidence interval is $\hat{\beta}_8 \pm z_{\alpha/2} \hat{se} = [0.310, 0.326]$.